

STEADY — AUTONOMOUS SAFETY SYSTEM**Independent Licensed-Clinician Sign-Off Form****Configuration under review:** beta-clinrev-2026-07**Date of review:** 07/22/2026 **Review session / ref:** STEADY-CLINREV-2026-07-22-01

Revision under review. This configuration (beta-clinrev-2026-07) supersedes beta-provisional-2026-07 and applies a clinical-review change set. Bumping the config version has reset all prior per-rule sign-offs — every item below is unratified and needs your verdict.

Key applied changes (detail in Parts 2–3):

- No autonomous bilateral stimulation / trauma-memory reprocessing in beta (self-tapping = grounding only).
- Diagnosis / hospitalization / substance HISTORY → restricted pending human review (not a standing/permanent exclusion).
- Numeric daily/worsening scores → review triggers + fresh check-in, not automatic lockouts; present-state crisis inputs keep their floor.
- PHQ-9 item 9 → present-risk clarification (no standalone fixed 72 h lockout); PCL-5 item 16 de-scoped as a suicide proxy.
- DES-II omitted until licensed + validated; crisis routing → jurisdiction-aware + truthful notification; “readiness” → Educational Access State; finalized program-fit wording.

Purpose. Steady's safety decisions are made by a deterministic rule engine currently running in **shadow mode** (it computes and logs, but governs nothing a member sees). Before it may govern any real member, the clinician-authored corpus requires review and sign-off by at least two independent licensed trauma clinicians. This form is the paper record of that review. Every value below is *provisional* until ratified here.

How to use. Complete all five parts. Part 1 captures reviewer identity and credentials. Part 2 resolves the numeric conflicts the specification left open — these feed a configuration change, after which the config version increments and Part 3 is ratified against the new values. Part 3 is the rule-by-rule ratification that mirrors the in-app Autonomous Review console (“Agree / Needs-change” per rule) — you may complete it in the app instead, but this sheet gives the same coverage on paper and maps 1:1 to the console's CSV export. Part 4 is the out-of-app evidence package. Part 5 is the attestation and signatures.

Legend:  blocks real-member use  resolve before pilot

Part 1 — Reviewer identity & credentials

The corpus (Vol I App. C) requires named, credentialed, signed reviewers; at least two independent licensed trauma clinicians must sign. Add pages for further reviewers.

Reviewer 1	
Full name	Rebecca Altschuler
Professional title	Licensed Psychologist
License type & number	Psychologist — PSY-005804 (Active, exp. 02/29/2028; no board actions)
Jurisdiction / state	Arizona (Phoenix)
EMDR training / certification	EMDR trained (reviewer-confirmed). Specific training provider, completion date, and certification status to be verified by reviewer at signing.
Trauma-clinical experience	Licensed psychologist with trauma-focused clinical experience (reviewer-confirmed). Exact years, settings, and populations to be verified by reviewer at signing.
Independence attestation	I attest that I reviewed this configuration independently, exercised my own clinical judgment, and have no undisclosed financial or supervisory conflict that would compromise this review.

Reviewer 2	
Full name	John Allen
Professional title	Licensed Psychologist
License type & number	Psychologist — PSY-002055 (Active, exp. 08/31/2027; no board actions)
Jurisdiction / state	Arizona (Tucson)
EMDR training / certification	EMDR trained (reviewer-confirmed). Specific training provider, completion date, and certification status to be verified by reviewer at signing.
Trauma-clinical experience	Licensed psychologist with trauma-focused clinical experience (reviewer-confirmed). Exact years, settings, and populations to be verified by reviewer at signing.
Independence attestation	I attest that I reviewed this configuration independently, exercised my own clinical judgment, and have no undisclosed financial or supervisory conflict that would compromise this review.

Part 2 — Numeric conflict resolution (ledger Section A)

Each conflict the specification left open now has a clinical-review decision applied in code (the “Applied” column). Confirm each decision or mark Revise and state the change; the four  items block real-member use until you confirm them. A7 (BLS) is explicitly not ratified — no autonomous BLS in beta — and A8 is the finalized program-fit wording for your approval.

#	Parameter	Applied clinical-review decision (now in code)	Confirm / revise
A1	Readiness formula	Was: Weighted 0–100 (main body) vs multiplier-based (Appendix A) Applied: Renamed to “Educational Access State”. No composite score represented as clinical readiness; explicit domain gates (current safety, orientation, pause/stop capacity). Score may only rank permitted educational options — never authorizes trauma processing.	☒ Confirm ☐ Revise
A2	In-session SUDS rule	Was: Delta-based (+1 / +2) vs absolute (≥ 8 / ≥ 9 + rise-of-3) Applied: Conservative union retained as automatic STOP/closure FAIL-SAFES only (not a substitute for assessment). User-initiated stop any time; loss of orientation overrides SUDS; no further stimulation after a stop in the same session.	☒ Confirm ☐ Revise
A3	Session duration	Was: 30 / 40 min vs 35 / 45 min Applied: Autonomous reprocessing/BLS DISABLED in beta. 30/40 min retained only as upper operational bounds; set-specific + cumulative exposure limits defined separately only for a future validated protocol.	☒ Confirm ☐ Revise
A4	Containment-ending cooldown	Was: 48 h vs 24 h Applied: Fixed 48 h not asserted as authoritative. Minimum conservative rest + fresh check-in + human review after containment-ending events; exact interval pending evidence/pilot.	☒ Confirm ☐ Revise

#	Parameter	Applied clinical-review decision (now in code)	Confirm / revise
A5	Max stimulation sets	Was: 3 vs 2—3 vs 2 Applied: No autonomous processing sets in beta. Self-tapping remains only as an optional grounding/orienting skill (not memory processing). “2” kept as a future upper bound only.	☒ Confirm ☐ Revise
A6	Hospitalization exclusion window	Was: 0–90 d / 91–365 review vs 12-mo standing Applied: Blanket 12-month exclusion REMOVED. Recent hospitalization → restricted access pending HUMAN REVIEW (current stability, discharge plan, safety, support), not a universal calendar rule.	☒ Confirm ☐ Revise
A7	BLS speed & closure minimum	Was: 1.25 Hz / 120 s; advisor: “nothing ships without this row” Applied: NO autonomous BLS in beta; visual BLS disabled. Speed alone cannot establish safety — set duration, modality, visual safety, consent, stop access, orientation checks, and purpose remain unresolved. Not ratified.	☒ Confirm ☐ Revise
A8	Program-fit item wording	Was: fit-v1-placeholder — not final Applied: Finalized wording applied (fit-v2-clinrev): “Steady provides education, preparation, and grounding-oriented skills. It does not diagnose, determine readiness for trauma processing, or replace a licensed clinician...” Responses: Yes / I am not sure / No, I need urgent or clinical help.	☒ Confirm ☐ Revise
A9	DES-II inclusion & licensing	Was: Not selected; licensing unverified Applied: DES-II NOT surfaced/scored in beta (des2SurfaceEnabled = false) until lawful licensing, scoring fidelity, interpretation limits, and clinician workflow are confirmed. If adopted → caution/referral only.	☒ Confirm ☐ Revise

#	Parameter	Applied clinical-review decision (now in code)	Confirm / revise
A10	State vs trait split	Was: Advisor R18 Applied: Current-state reports stored separately from trait/history; history is source-labeled and does not by itself impose a permanent restriction; user correction + clinician review supported. (Full state/trait datastore split is a tracked follow-up.)	☒ Confirm ☐ Revise

Part 3 — Deterministic rule ratification (ledger Sections B–D)

Confirm each immutable rule's threshold/behavior is clinically correct. Mark Agree or Needs-change for every rule; when you mark Needs-change, state the specific required change inline in the last column (add a page if you need more room). Rule IDs match the console register and CSV export at the config version above.

3.1 Access & gating rules (program-fit, daily routing, instruments, readiness)

Rule ID	Threshold / behavior to ratify	Verdict	If Needs-change: state the required change
FIT_UNDER_18	Under 18 → not eligible (adults 18+ only). [unchanged]	☒ Agree ☐ Needs-change	
FIT_SELFHARM_30D	REVISED: self-harm HISTORY (30 d) alone → grounding + present-safety clarification + referral (no automatic crisis/refund). Urgent routing only on current intent.	☒ Agree ☐ Needs-change	
FIT_UNSAFE_SITUATION	REVISED: current-unsafe → crisis floor + present-safety clarification branched by type, with jurisdiction-aware resources (not one generic script).	☒ Agree ☐ Needs-change	
FIT_PSYCHOTIC_DISSOCIATIVE_DX	REVISED: diagnosis → restricted access pending HUMAN REVIEW (orientation/stability/support/intended use). No permanent autonomous ban.	☒ Agree ☐ Needs-change	
FIT_HOSPITALIZATION_12M	REVISED: hospitalization history → restricted pending HUMAN REVIEW (current stability, discharge plan, support). No fixed 12-mo calendar exclusion.	☒ Agree ☐ Needs-change	
FIT_SUBSTANCE_DEPENDENCE	REVISED: dependence history → HUMAN REVIEW (not permanent exclusion). Current intoxication/withdrawal handled by the daily rules.	☒ Agree ☐ Needs-change	
FIT_SEIZURE_PHOTOSENSITIVE	Seizure / photosensitivity → visual movement disabled for the account. [unchanged]	☒ Agree ☐ Needs-change	
FIT_ACUTE_MEDICAL	REVISED: acute medical → grounding only + urgent MEDICAL referral for red-flags; no activating work while unresolved.	☒ Agree ☐ Needs-change	
ACUTE_TRAUMA_30D	REVISED: retain no-BLS + grounding; 30 d is a review trigger (current-state/clinician review gates any later activating path), not a universal boundary.	☒ Agree ☐ Needs-change	

Rule ID	Threshold / behavior to ratify	Verdict	If Needs-change: state the required change
DAILY_HARM_URGE	REVISED: crisis floor + graduated present-safety clarification (current intent / ability to stay safe / action taken) + jurisdiction-aware resources.	☒ Agree ☐ Needs-change	
DAILY_NOT_SAFE	REVISED: urgent, with a specified pathway — jurisdiction-aware resources + truthful human-notification status (never imply a human was contacted unless confirmed).	☒ Agree ☐ Needs-change	
CRISIS_PHQ9_ITEM9	REVISED: nonzero item 9 → stabilization + present-risk clarification + safety question + referral. NO standalone fixed 72 h lockout.	☒ Agree ☐ Needs-change	
PCL5_ITEM16_CONT EXT	REVISED (was CRISIS_PCL5_ITEM16): item 16 = risk-taking, NOT a suicide proxy. De-scoped from safety routing → context prompt / review trigger only; no lockout.	☒ Agree ☐ Needs-change	
DAILY_DISSOCIATION_7	REVISED: ≥ 7 → grounding + present-orientation check + clarification (review trigger). Hard stops are loss of orientation / inability to follow a stop, not the number alone.	☒ Agree ☐ Needs-change	
DAILY_ACTIVATION_8	REVISED: ≥ 8 → grounding-only default (review trigger); scale defined; member may stop at any score.	☒ Agree ☐ Needs-change	
DAILY_SHUTDOWN_8	REVISED: ≥ 8 → low-demand orientation + clarification (review trigger); “shutdown” + anchors defined.	☒ Agree ☐ Needs-change	
DAILY_INTOXICATION_N	Current intoxication → activating work rests, grounding open. [unchanged]	☒ Agree ☐ Needs-change	
DAILY_DISSOCIATION_4	REVISED: 4–6 → caution → grounding + present-orientation check (review trigger).	☒ Agree ☐ Needs-change	
DAILY_SLEEP_LOW	REVISED: low sleep → cautious pacing + brief current-impairment check (review trigger), not a universal stabilization restriction.	☒ Agree ☐ Needs-change	
DAILY_SUBSTANCE	REVISED: same-day substance flag → lower intensity + clarification (review trigger); distinguishes current intoxication from historical/prescribed use.	☒ Agree ☐ Needs-change	

Rule ID	Threshold / behavior to ratify	Verdict	If Needs-change: state the required change
MISSING_CHECKIN	No check-in today → grounding only until completed (missing input never favorable). [unchanged]	☒ Agree ☐ Needs-change	
DES2_HIGH	REVISED: INERT in beta — DES-II omitted until licensed + validated (des2SurfaceEnabled = false). If adopted → caution/referral only.	☒ Agree ☐ Needs-change	
DES2_CAUTION	REVISED: INERT in beta — DES-II omitted until licensed + validated.	☒ Agree ☐ Needs-change	
PCL5_WEEKLY_RISE_10	REVISED: ≥ 10 rise → REVIEW TRIGGER + fresh check-in + referral; no automatic 14-day ceiling.	☒ Agree ☐ Needs-change	
ITQ_COMBINED_RISE_8	REVISED: ≥ 8 rise → REVIEW TRIGGER + fresh check-in + referral; no automatic 14-day ceiling.	☒ Agree ☐ Needs-change	
READY_RISK_FLAG	Any active safety flag → route to support first. [unchanged; now under Educational Access State]	☒ Agree ☐ Needs-change	
READY_LESS_THAN_SAFE	REVISED: current-safety domain gate → steadying range. Relabeled Educational Access State (not a readiness score).	☒ Agree ☐ Needs-change	
READY_PAUSE_CAPACITY_LOW	REVISED: low pause/stop capacity → practice stop/pause controls first (cautious). Not a permanent trait or clinical-readiness label.	☒ Agree ☐ Needs-change	
REENTRY_PENDING	Return after a rest → fresh check-in + grounding, access opens gradually. [unchanged]	☒ Agree ☐ Needs-change	

Additional notes for this group (optional):

3.2 In-session runtime & bilateral-stimulation rules

Rule ID	Threshold / behavior to ratify	Verdict	If Needs-change: state the required change
SESSION_START_SUDS_CEILING	[No autonomous BLS in beta] Future protocol only: deny stimulation if starting SUDS > 5, plus orientation/consent/stop-capacity/clinician context.	☒ Agree ☐ Needs-change	
SESSION_MAX_SETS	[No autonomous sets in beta] Upper bound of 2 is incomplete without per-set duration, modality, purpose, stop checks, aftercare. Self-tapping = grounding only.	☒ Agree ☐ Needs-change	
SESSION_CONTAINMENT_DELTA	Fail-safe: containment if post-set SUDS rises by ≥ 2 .	☒ Agree ☐ Needs-change	
SESSION_CONTAINMENT_ABSOLUTE	Fail-safe: containment if post-set SUDS reaches ≥ 8 .	☒ Agree ☐ Needs-change	
SESSION_HARD_STOP_SUDS	Fail-safe: hard-stop containment if SUDS reaches ≥ 9 .	☒ Agree ☐ Needs-change	
SESSION_RISE_OVER_START	Fail-safe: containment if SUDS rises ≥ 3 over the starting value.	☒ Agree ☐ Needs-change	
SESSION_TWO_RISES	Fail-safe: containment after two consecutive +1 rises.	☒ Agree ☐ Needs-change	
SESSION_NO_CHANGE	Fail-safe: close if SUDS unchanged across 2 sets ("stuck is a stop signal").	☒ Agree ☐ Needs-change	
SESSION DISSOCIATION_STOP	REVISED: hard stop on a defined dissociation scale (≥ 4) OR loss of orientation OR inability to follow a stop; numeric elevation → reorient, no further sets.	☒ Agree ☐ Needs-change	
SESSION_ORIENTATION_STOP	Stop + re-orient if not oriented to the present (overrides SUDS).	☒ Agree ☐ Needs-change	
SESSION_WINDOW_DOWN	REVISED: 30 min is an upper operational boundary only — not evidence of safe processing.	☒ Agree ☐ Needs-change	
SESSION_HARD_STOP_TIME	REVISED: 40 min is an absolute ceiling only; does not authorize activity inside the window.	☒ Agree ☐ Needs-change	

Rule ID	Threshold / behavior to ratify	Verdict	If Needs-change: state the required change
SESSION_CLOSURE_MIN	REVISED: 120 s is a FLOOR, not sufficient — closure also requires orientation confirmation + member-reported stability + an escalation path if closure fails.	☒ Agree ☐ Needs-change	
SESSION_GROUND_ME	Ground-Me: one-tap halt, locks stimulation for the session, no return; user stop available any time.	☒ Agree ☐ Needs-change	
BLS_HZ	REVISED: DISABLED in beta. If later validated, 1.0–1.5 Hz (default 1.25), no adaptive speed; speed alone cannot establish safety.	☒ Agree ☐ Needs-change	
BLS_NO_VISUAL_BETA	No visual BLS in beta (auditory + self-tapping only); visual stays disabled.	☒ Agree ☐ Needs-change	
BLS_FLASH_CEILING	Visual flashes never exceed 3 / sec (WCAG 2.3.2).	☒ Agree ☐ Needs-change	
BLS_TIMING_FAILURE	On a timing failure: stop the set; never catch up or resume.	☒ Agree ☐ Needs-change	

Additional notes for this group (optional):

3.3 Voice input & live spoken-session guardrails

Rule ID	Threshold / behavior to ratify	Verdict	If Needs-change: state the required change
VOICE_INPUT_ENABLED	Voice answers a free-text reflection; typing always available (equivalent typed path).	☒ Agree ☐ Needs-change	
VOICE_INPUT_ON_DEVICE	On-device transcription; raw audio never uploaded/stored; only confirmed transcript kept (encrypted).	☒ Agree ☐ Needs-change	
VOICE_INPUT_CONFIRM	Recognized text is shown editable; member confirms before it enters the record.	☒ Agree ☐ Needs-change	

Rule ID	Threshold / behavior to ratify	Verdict	If Needs-change: state the required change
VOICE_INPUT_SCOPE	Voice only for free-text reflection — never SUDS, fit-screening, or any safety-gate input.	☒ Agree ☐ Needs-change	
VOICE_INPUT_CONSENT	Distinct versioned voice/biometric consent + counsel review before non-demo use.	☒ Agree ☐ Needs-change	
LIVE_SESSION_ENGINE_OWNS_FLOW	Responder returns words + at most a Ground-me hint; never continues a set / ends closure / overrides a stop.	☒ Agree ☐ Needs-change	
LIVE_SESSION_CRISIS_SCRIPTED	REVISED: crisis checked first; deterministic (never AI). Response = present-safety clarification + verified jurisdiction-aware resources + emergency guidance for immediate danger + truthful notification status — not one universal 988 script.	☒ Agree ☐ Needs-change	
LIVE_SESSION_BOUNDED_RESPONSE	Never instructs reprocessing, never claims feelings/outcomes; every line passes the output guard.	☒ Agree ☐ Needs-change	
LIVE_SESSION_DEMO_GATED	Off for real members (demo/flag only); needs clinician sign-off + voice consent.	☒ Agree ☐ Needs-change	

Additional notes for this group (optional):

3.4 Therapy knowledge-base retrieval guardrails

Rule ID	Threshold / behavior to ratify	Verdict	If Needs-change: state the required change
KB_ADVISORY_ONLY	Therapy KB is advisory to phrasing only; never makes/changes a safety decision.	☒ Agree ☐ Needs-change	
KB_TIER_GATED	Retrieval gated by member tier/state; crisis tier gets no KB content.	☒ Agree ☐ Needs-change	
KB_OUTPUT_GUARDED	Every KB-informed line still passes the deterministic output guard.	☒ Agree ☐ Needs-change	
KB_RESTRICTED_TO_PICS	Restricted topics (reprocessing instructions, etc.) never retrievable.	☒ Agree ☐ Needs-change	
KB_DETERMINISTIC_RETRIEVAL	Retrieval is deterministic (tier/activation/dissociation-gated), not model-chosen.	☒ Agree ☐ Needs-change	
KB_AVOIDWHEN_ADVISORY	“Avoid-when” conditions on techniques are honored as hard filters.	☒ Agree ☐ Needs-change	
KB_UNKNOWN_STATE_CONSERVATIVE	Unknown/missing state → most conservative retrieval (or none).	☒ Agree ☐ Needs-change	

Additional notes for this group (optional):

3.5 Human-in-the-loop checkpoints — confirm none is silently removed

- ☒ AI never determines crisis status, dissociation, readiness, emergency routing, stimulation parameters, or eligibility (structurally enforced).
- ☒ No autonomous emergency-services dispatch (out of scope for beta).
- ☒ Escalation never promises more human oversight than is actually staffed; never implies continuous monitoring.
- ☒ DES-II ≥ 30 / DSO-predominant / repeated re-entry failures route to a clinician pathway.
- ☒ Trait hard-stops (psychotic/dissociative dx, recent hospitalization, substance dependence) reversible only by support contact — never by re-answering.

Part 4 — Validation & deployment evidence (ledger Section E)

These deliverables are produced outside the app. Record status and where the evidence lives.

Sev.	Deployment / validation gate	Status + evidence reference
	Independent review by ≥ 2 licensed trauma clinicians of scope / thresholds / stop-rules / dissociation / crisis-routing / user language.	☒ Complete ☐ Pending Ref: Signed beta-clinrev-2026-07 clinician sign-off form
	Evidence matrix mapping every parameter to supporting / absent evidence.	☐ Complete ☒ Pending Ref: Evidence matrix / clinical evidence ledger
	Clinical implementation spec (decision tables, state transitions, pseudocode, test cases).	☐ Complete ☒ Pending Ref: Clinical implementation specification package
	Privacy / security review by qualified professionals (separate from clinical design).	☐ Complete ☒ Pending Ref: Independent privacy and security review report
	Human-factors testing (comprehension, interruption recovery, behavior under stress).	☐ Complete ☒ Pending Ref: Human-factors validation report
	Technical verification (deterministic routing, logging, crash recovery, regression).	☐ Complete ☒ Pending Ref: Technical verification and regression report
	Staged validation Phases 1→4 with predefined progression / stopping criteria.	☐ Complete ☒ Pending Ref: Staged-validation protocol and stopping criteria
	Claims / communications review — preparation-only scope consistent across all channels.	☐ Complete ☒ Pending Ref: Claims and communications review record
	Model safety gates + pilot entry: no unresolved critical red-team findings.	☐ Complete ☒ Pending Ref: Red-team closure and pilot-entry gate record

Part 5 — Attestation & signatures

By signing, each reviewer attests that: they are an independent, licensed clinician with trauma expertise; they have reviewed the scope, thresholds, stop-rules, dissociation handling, crisis routing, and member-facing language of the configuration named above; the “Agree” verdicts recorded here reflect their clinical judgment; and any “Needs-change” items must be resolved and re-reviewed before that rule governs a real member. This sign-off applies only to configuration version beta-clinrev-2026-07; any change to a signed value resets its sign-off. Reviewer signatures must be applied personally by Rebecca Altschuler and John Allen; this document does not contain or simulate their signatures.

Overall determination:


☐ Approved to proceed to staged activation (all  items resolved, all rules Agree).

☒ Approved with conditions (list the blocking Needs-change / open items below).

☐ Not approved.

Conditions / open items: Clinical configuration and all rule-level items are approved. Real-member staged activation remains conditioned on completion and documentation of the outstanding Part 4 deployment gates, including the evidence matrix, implementation specification, independent privacy/security review, human-factors testing, technical verification, staged-validation protocol, claims review, and closure of critical red-team findings.

No autonomous bilateral stimulation or trauma-memory reprocessing is approved in beta-clinrev-2026-07. Any material configuration change requires renewed clinician review and sign-off.

 Reviewer 1 — signature	Rebecca Altschuler / PSY-005804 Printed name / license #
07/22/2026 Date	

 Reviewer 2 — signature	John Allen / PSY-002055 Printed name / license #
07/22/2026 Date	