

Steady — Presentation Layer Handoff

How data, gating, routing, and processing should be rendered

Version 1.0 · First pass · Prepared for engineering

0. Read this first

This handoff has one governing finding, and it changes the shape of the problem.

Your clinical spec has already decided the answer to “how should we show the data.”

Volume 2, Part 1 and the Row 39–53 registry state it repeatedly and without ambiguity:

Scores, diagnostic bands, criteria labels, and hidden track names are not displayed to members... The member experiences routing as which practices appear today. If any surface in the product contradicts this paragraph, it is a defect.

And on the readiness score specifically: *“keep score hidden”* appears on every single readiness row.

So the brief is not “design a dashboard.” The brief is the opposite: **build a translation layer that converts a dense deterministic routing engine into a surface that displays no numbers at all.** The data-orientation you’re seeing in the current build isn’t a styling problem — it’s an unbuilt abstraction. The engine’s internal state is leaking directly onto the screen because nothing has been written to stand between them.

Everything below is that layer.

One caveat on scope: I could not review your current screens. Both uploaded volumes are text-only clinical documents with no screenshots or wireframes. So this is a *target-state* spec derived from your clinical constraints plus competitive research — not a before/after audit. When you send screens or a Figma link, the gap analysis is a separate short pass.

1. Competitive scan — six platforms, GUI only

Researched from public sources: product pages, app store listings, design case studies, and published research. Where I’m inferring rather than observing, I say so.

1.1 NeuroFlow (Onward)

What it is visually: A two-sided system. The member app is a light, card-based wellness app — mood and sleep logging, a browsable content gallery, assigned activities, journals, breathing tools, assessments delivered as periodic check-ins, and a rewards/points system. The provider side is a population dashboard built on their proprietary NeuroFlow Severity Score, with risk heatmaps and cohort stratification by acuity.

The pattern worth stealing: *Assignment.* Content arrives as an assignment from a care team or is auto-assigned by severity tier, so the member is never staring at a library trying to self-diagnose. The routing does the choosing.

The pattern worth rejecting: *Points and rewards.* NeuroFlow gamifies with redeemable points and gift cards, and Penn's digital mental health group describes the behavioral-economics scaffolding as central to the design. This is directly incompatible with Steady. In a trauma-preparation context, extrinsic reward pressure competes with the clinical requirement that stopping be a valid choice (Vol 1, B-4). Do not build points.

The tension you must resolve: NeuroFlow shows severity scores — to providers, and trend data to members. Your spec forbids both on the member side. Since your memory-of-record has Steady repositioning toward enterprise buyers (payers, health systems, risk-bearing primary care), you are going to be asked for exactly the dashboard your clinical spec prohibits on the member surface. Section 3 resolves this with a hard audience split. Get this right early; retrofitting it is expensive.

1.2 Headspace (Ebb)

The most directly instructive case, because Headspace published their design process.

Their internal conversational-AI guidelines, visible in the Figma case study, were: always make the AI visible; differentiate from human-delivered care; create a reflective and contemplative space; reinforce privacy and safety at all times; honor member choice and autonomy. Their UX lead states plainly that they didn't want the AI to be invisible — members should know they're interacting with AI, not a human. They emphasize clinical-expert authorship early in the experience specifically because that's what moved trust in testing.

Steal: the trust sequencing. Clinical provenance surfaces *early*, not buried in a settings page. You have two ratifying psychologists and a 30-rule engine — that is a trust asset and it is currently invisible.

Steal: conversation deletion and exit as first-class controls, available at any point.

Reject: the mascot. Ebb is a friendly gradient blob; Wysa is a penguin; Finch is a bird you care for. Steady should not have a character. A mascot creates a relational object, and

relational objects invite attachment and dependency — a real risk in a C-PTSD population using an autonomous product with no clinician in the loop. Your differentiation is that Steady is an *instrument*, not a friend. Design it that way. This is a genuine fork in the road and I'd argue it hard.

1.3 Ash (Slingshot AI)

Voice-first, conversation-primary, with cross-session memory and a "weekly insight" surface. Their recent release notes mention making voice selection clearer and giving the weekly insight screen a calmer, cleaner look — a signal that their insight surface was over-dense and got walked back.

Steal: the weekly insight as a *narrative* artifact rather than a chart. Note that even Ash — with no equivalent of your no-numbers constraint — is trending toward calmer and less dense.

Watch: Slingshot markets Ash as built for therapy while the app itself disclaims being therapy. That gap is a live regulatory exposure in the category. Your wellness-lane copy discipline and CI-enforced banned vocabulary are the correct posture; don't drift toward Ash's marketing register.

1.4 Wysa

Conversational-everything: goal-setting, planning, and tools all live inside the chat UI rather than in separate sections. The self-care library uses clean cards with clear iconography. Notable pattern — the app passively compiles a journey log by summarizing completed exercises, so the member gets a record without being asked to produce one. Persistent SOS access. Palette is a deep near-black, a pale ice blue, and a mid teal.

Steal: *zero-effort history*. The member never fills out a log; the system assembles one from what they actually did. This is exactly right for Steady's referral export — the artifact a member hands a clinician should build itself.

Note: Wysa and Youper are text-only with limited cross-session memory, while Ash, Headspace Ebb, and others now offer voice plus persistent memory. Voice-first with memory is becoming table stakes in this category.

1.5 Calm

The reference case for sensory immersion — ambient scene backgrounds, mood check-ins that feed content recommendation, a large library handled with clean cards, muted color, and very low visual noise. Widely cited for a soft, non-aggressive payroll.

Steal: the ambient-scene approach to reducing arousal before any content loads, and the

restrained subscription surface. You have a three-tier model; Calm’s paywall is the benchmark for not making that feel predatory.

Reject: the browse-first library model. Choice paralysis is a documented failure mode for depleted users, and your routing engine exists precisely so members don’t have to choose.

1.6 Spring Health / Lyra (enterprise comparators)

Assessment-first flows. Spring Health runs a dynamic assessment screening for 12+ conditions, then produces a care plan and provider matches from the results. Lyra measures progress via PHQ-9 and GAD-7. One reviewer notes Lyra’s interface feels less visually engaging than Headspace’s — a useful warning that clinical rigor plus enterprise buyers tends to produce cold, form-heavy UI.

Steal: the care-plan-as-output pattern. The assessment doesn’t dead-end in a score; it produces a plan.

Reject: assessment-first as the opening experience. One reviewer flags the assessment-first flow as friction when you already know what you want. Your 14-step gate is longer than Spring’s. Section 5 addresses this directly — it’s your single biggest UX risk.

1.7 Summary table

Platform	Data display to user	Content selection	Persona	Key transferable pattern
NeuroFlow	Scores + trends (both sides)	Assigned by severity	None	Assignment over browsing
Headspace Ebb	Minimal, conversational	AI-recommended	Blob character	Early clinical-provenance trust
Ash	Weekly insight, being simplified	Conversation-led	Named voice	Insight as narrative
Wysa	Low; passive journey log	In-chat	Penguin	Self-assembling history
Calm	Mood check-in only	Browse + recommend	None	Sensory pre-load, soft paywall
Spring/Lyra	Assessment results, clinical	Care plan from assessment	None	Assessment produces a plan
Steady	None. Zero	Deterministic	None —	

(target)	numbers.	routing	instrument	—
----------	----------	---------	------------	---

Steady is the only product in this set that shows the member nothing. That is a constraint, and it is also the most distinctive thing about the product. Design into it rather than around it.

2. The three failure modes to design against

Leakage. Any screen that shows a score, band, track name, criteria label, or trend chart is a defect per Vol 2. This includes soft leaks: a progress ring that implies a percentage, a color that maps to severity, a badge that reveals track membership, a chart of check-ins over time.

Denial framed as failure. Your engine narrows a member’s day constantly — cooldowns, caps, stabilization holds, imagery restrictions. If narrowing reads as “you failed the check,” you produce shame in a population where shame drives disengagement. Vol 2 is explicit: a narrow day is framed as *a grounding day, not a failure*.

Density under load. Vol 1 B-6 requires one primary task per screen, minimal reading during activation, and calm predictable visuals — because trauma symptoms impair concentration and working memory. A screen that is fine for you at your desk may be unusable for a member at 2am.

3. Audience architecture — the split that must happen now

Your prompt said “no matter who the user is.” There are four, and they need genuinely different surfaces off one data core.

Surface	Sees	Never sees	Purpose
Member	Today’s available practices, session flow, crisis access, self-assembled history	Any number, score, band, label, track, chart	Do the practice
Referral export (member-controlled)	Structured summary the member hands a clinician; includes high-intensity trigger items excluded from self-guided use	—	Bridge to human care

Clinician / supervised mode	Instrument scores, routing decisions, rule fires, session logs	Raw free-text beyond minimum necessary	Clinical oversight
Enterprise / admin	Cohort-level aggregates, engagement, safety-event rates	Anything individually identifying	Buyer reporting

Engineering requirement: these are not view permissions on one component tree. Build the member renderer so it is *structurally incapable* of receiving a score. The API that serves the member surface should not return score fields at all — not hidden, not null, not filtered client-side. If a score never crosses the boundary, leakage becomes impossible rather than merely prohibited. Add a CI check on the member API contract in the same spirit as your existing banned-vocabulary check.

4. The member surface — Day States

Collapse the entire routing hierarchy into a small, fixed set of member-visible day shapes. The engine can evaluate 30 rules across six domains; the member sees one of five days.

```
DayState {
  shape:      OPEN | NARROW | STABILIZING | PAUSED | CRISIS
  practices:  [ practice refs available today ]
  primary:    practice ref      // the single recommended next action
  message_key: copy key        // governed copy, versioned
  human_support: always present
  // NO scores. NO bands. NO track. NO rule IDs.
}
```

OPEN — full set available. Primary action offered, everything else secondary. **NARROW** — grounding and stabilization practices only. Copy frames this as the day's work, not as a restriction. No explanation of *why* beyond a plain, non-clinical sentence. **STABILIZING** — a 72-hour hold is active. Fixed safety options surface. The day still has content; it is not an empty screen. **PAUSED** — a standing restriction applies. Referral surface is prominent. This is the state most likely to produce abandonment, so it needs the most careful copy. **CRISIS** — takes over the entire viewport. All other navigation is removed. Single-purpose, high-legibility, no branding flourish.

Rendering rule for narrowing: the interface should not show the member what they *would* have had. No greyed-out cards, no locked icons, no "unavailable" states with padlocks.

Absent is absent. A locked card is a scoreboard of what you failed to unlock.

This is the single most important line in this document.

5. The gate — 14 steps is the biggest UX risk in the product

Fourteen steps is long. Spring Health's assessment-first flow already draws friction complaints, and Steady's is longer and more sensitive.

Do not render this as a form. Render it as a paced sequence with these properties:

- **One question per screen.** Vol 1 B-6 requires it.
- **Resumable by default, saved continuously.** Vol 1's 2026 guidance and the "depleted brain" principle both require that pausing mid-check loses nothing. Never a "start over."
- **Progress shown as position, not percentage.** A simple sequence marker. Avoid a progress bar reading 30% — percentage framing invites abandonment maths.
- **Safety items commit immediately.** Your own pseudocode already specifies this: persist the response and evaluate the fixed rule the moment a safety-relevant item is answered, without waiting for submission. The UI must respect this — a safety disposition that has fired cannot be undone by backing out of the questionnaire.
- **No result screen with a number.** The gate terminates in a Day State, not a score. This is the care-plan-as-output pattern from Spring, minus the score.
- **Denial is not an error state.** If the gate routes to PAUSED or STABILIZING, that is a legitimate destination rendered in the same visual register as any other. No red. No warning iconography. No apology.

Exit affordance on every step. Labeled as a pause, not a quit — the "no-guilt close." Consistent position across all 14.

6. Session flow — rendering the state machine

Your Vol 2 state machine maps almost one-to-one onto screens:

```
created → pre_session_check → authorized → set_active → post_set_reassessment  
        → set_active (only after explicit authorization) → closure → completed
```

Design notes per state:

pre_session_check — brief, calm, no new information. If it denies, transition to the denial screen in the same visual register (see above).

authorized → set_active — this transition needs an explicit confirmation beat. Vol 1 B-6 calls for explicit transition confirmations. The member should have to actively step into the set, not slide into it.

set_active — the quietest screen in the product. BLS at a fixed 1.25 Hz default, range 1.0–1.5. **No adaptive speed.** Stop and pause must be permanently visible, high-contrast, and reachable without precision — full-width tap targets, not small icons. Grounding (“ground me”) is a peer action to stop, not buried.

post_set_reassessment — the highest-consequence input in the flow. Minimal reading. Single scale. Do not show the member their previous answer for comparison; that invites performance.

closure — mandatory. Should feel like an arrival, not a form to clear. This is the screen most likely to be skipped and most clinically important not to skip.

technical_interruption → recovery — design this properly rather than treating it as an edge case. A member dropping mid-set on a bad connection needs a defined re-entry that respects any safety state that fired before the drop.

Any state → crisis_support — reachable from every screen at a fixed position, per the persistent-crisis-affordance standard. Fixed position is a safety requirement, not a layout preference: it must be findable without reading.

7. Visual system

Directional, not final. Justified against the brief rather than pulled from a wellness template.

Palette

The category default is de-saturated teal and sage. It’s fine, it’s also what everyone uses, and it reads as generic wellness. Steady’s own vernacular is different: steadiness, level, horizon, ballast — instrument language, not spa language. Cool, low-chroma, quiet.

```
--steady-deep      #1C222C    // ground, session backdrop
--steady-slate     #333C4A    // primary text on light
--steady-horizon   #5B7A8C    // primary interactive — dusty blue-grey
--steady-mist      #E8EBE9    // cool paper, primary light surface
--steady-moss      #7A8B6F    // available / affirmative, low saturation
--steady-amber     #C08A4E    // attention only — never alarm
```


No red or orange anywhere in the system, including destructive and crisis actions. High-vibrancy red is processed as threat, which is the wrong physiological response to trigger in this population — including on the crisis screen, where the member is already activated. Crisis uses `--steady-deep` with maximal contrast and typographic weight, not color alarm.

`--steady-amber` appears only for attention states and never for alarm; if it starts showing up on more than a couple of surfaces, that's a signal the design has drifted.

Type

One family, two optical roles. The conventional move is a display serif paired with a body sans; a single humanist sans with high x-height and open counters, differentiated by scale and tracking rather than by contrast, reads as one calm voice and avoids the two-voice tension. Vol 1 requires legibility under fatigue and cognitive load — high x-height and open counters are the specific properties that survive tired eyes.

Minimum body size 17px. Line length capped at ~60 characters. Line height generous — 1.6 minimum on body copy.

Signature element: the horizon

One thin horizontal rule, at a consistent position on every screen, that sits **lower on a narrow day and higher on an open day**. It carries the day's shape without a number, a color code, or a label. During a set, the BLS indicator travels along that same line — so the product's core mechanic and its state indicator are the same object.

Boundary condition, and this matters: a single static position per day is a state indicator, not a chart. The moment you animate it across days, show history, or let a member scrub back through it, it becomes a trend chart and violates Vol 2. Build it stateless. If someone later asks for "just a little history on the horizon," that's the defect Vol 2 is warning about.

Motion

Slow, few, and always interruptible. Respect `prefers-reduced-motion` as a hard requirement, not a nicety — you already remove visual BLS globally for photosensitivity and seizure risk, so motion sensitivity is an established first-class concern in this product.

8. Component contracts for engineering

Component	Receives	Must never receive
DayCanvas	DayState	scores, rule IDs, track name

PracticeCard	practice ref, availability	severity, reason-for-exclusion
GateStep	one item, position in sequence	running total, partial score
SessionSurface	fixed BLS params, controls	adaptive/derived speed
ClosureSurface	closure target	pre/post comparison
CrisisSurface	fixed dispositions only	anything model-generated
HistoryStrip	completed practice refs	dates-as-streak, counts, charts

Two hard engineering rules:

1. **Crisis content is never model-generated.** Vol 2 specifies fixed dispositions with versioned copy. The crisis surface renders governed copy keys — no LLM in that path, ever.
2. **Every member-facing string is a versioned copy key**, not an inline literal. Your CI already blocks banned vocabulary (treat, therapy, diagnose, reprocessing). Copy keys make that check enforceable at build time and make the Row 15 clinical copy review a diff instead of an archaeology project.

9. Sequencing

First — the boundary. Split the member API so scores cannot cross it. Everything else depends on this and it gets harder every week you wait.

Second — Day States. Collapse routing into the five shapes. This alone converts most of the current data-orientation into something a member can read.

Third — the gate. Rebuild as a paced, resumable sequence. Highest abandonment risk in the product.

Fourth — session and closure. Apply the visual system where it matters most.

Fifth — clinician and enterprise surfaces. These can be data-dense; that's what they're for. Building them last prevents their requirements from contaminating the member surface, which is the most common way products like this go wrong.

10. Open questions

1. Does the horizon element survive clinical review, or does any persistent state indicator read as a covert score?
2. Where does the AI companion live spatially? It isn't in the Vol 2 state machine at all — it needs a defined position relative to the routing engine before it gets a position on screen.
3. Does the referral export assemble passively (Wysa model) or require member curation?
4. Three tiers across five day shapes: what does a STABILIZING day look like on the entry tier versus the top tier? If tiers gate safety content, that's a clinical question before it's a design one.
5. What renders during PAUSED for a member under a 0–90 day exclusion? That member has weeks of days with no practice available and is the most likely to churn and the most in need of a bridge to human care.